

Machine Learning 1 Naïve Bayes




Types of Learning Problems

- **Supervised learning:** correct answers for each training example
 - **Classification:** learning predictor with *discrete* outputs
 - **Regression:** learning predictor with *real-valued* outputs
- **Unsupervised learning:** no correct answers, just find good representations / features of the data
- **Reinforcement learning:** reward function, no correct answers

Machine Learning Roadmap


- Define the problem
 - Type of problem, domain (i.e. spam filtering, digit recognition)
- Look at several learning approaches / models
 - *Naïve Bayes* (today), *Perceptrons*, *Logistic Regression* (next week)
- How to find model parameters:
 - Maximum Likelihood
- Themes throughout
 - Working with data
 - Preventing overfitting
 - Evaluating performance

Classification



Classification and Machine Learning


- Dataset: each data point, x , is associated with some **label** (aka class), y
- Goal of classification: given inputs x , write an algorithm to predict labels y
- Workflow of classification process:
 - Input is provided to you
 - Extract **features** from the input: attributes of the input that characterize each x and hopefully help with classification
 - Run some machine learning algorithm on the features: today, *Naïve Bayes*
 - Output a predicted label y



Training

- Big idea: ML algorithms learn patterns between features and labels from *data*

- You don't have to reason about the data yourself
- You're given **training data**: lots of example datapoints and their actual labels



Training: Learn patterns from labeled data, and periodically test how well you're doing



Eventually, use your algorithm to predict labels for unlabeled data

Example: Spam Filter

- Input: an email
- Output: spam/ham
- Setup:
 - Get a large collection of example emails, each labeled “spam” or “ham”
 - Note: someone has to hand label all this data!
 - Want to learn to predict labels of new, future emails
- Features: The attributes used to make the ham / spam decision
 - Words: FREE!
 - Text Patterns: \$dd, CAPS
 - Non-text: SenderInContacts, WidelyBroadcast
 - ...



Dear Sir.

First, I must solicit your confidence in this transaction, this is by virtue of its nature as being utterly confidential and top secret. ...

TO BE REMOVED FROM FUTURE MAILINGS, SIMPLY REPLY TO THIS MESSAGE AND PUT "REMOVE" IN THE SUBJECT.

99 MILLION EMAIL ADDRESSES
FOR ONLY \$99

Ok, I know this is blatantly OT but I'm beginning to go insane. Had an old Dell Dimension XPS sitting in the corner and decided to put it to use, I know it was working pre being stuck in the corner, but when I plugged it in, hit the power nothing happened.

Example: Digit Recognition

- Input: images / pixel grids


- Output: a digit 0-9

- Setup:

- Get a large collection of example images, each labeled with a digit
- Note: someone has to hand label all this data!
- Want to learn to predict labels of new, future digit images

- Features: The attributes used to make the digit decision

- Pixels: (1,1)=ON
- Shape Patterns: NumComponents (Number of connected regions), AspectRatio (Height-to-width ratio), NumLoops (Number of closed loops (e.g., '8' has two loops)).
- These features help the model distinguish between similar-looking digits like '1' and '7' or '3' and '8'.



0

1

2

1

??

Other Classification Tasks

- Classification: given inputs x , predict labels (classes) y


- Examples:

- Object recognition
input: images; classes: object type
- Medical diagnosis
input: symptoms; classes: diseases
- Automatic essay grading
input: document; classes: grades
- Fraud detection
input: account activity; classes: fraud / no fraud
- Customer service email routing
input: a customer service email; classes: billing, technical support, account management, sales, shipping, returns, product inquiry, etc.
- ... many more

- Classification is an important commercial technology!



Model-Based Classification



Model-Based Classification

■ Model-based approach

- Build a model (e.g. [Naïve Bayes](#)) where *both the label and features are random variables*
- Instantiate any observed features
- Query for the distribution of the label conditioned on the features

■ Challenges

- What structure should the model have?
- How should we learn its parameters?



Naïve Bayes Model

- Naïve Bayes: Assume all features are independent effects of the label.


- This simplification is the “**naïve**” **assumption**: it ignores correlations between features, making computation easier
- While this assumption is often not true in reality, it simplifies computations and still performs surprisingly well in practice.

- Random variables:

- Y = The label
- F_1, F_2, \dots, F_n = The n features

- Parameters: Probability tables:

- $P(Y)$ = Probability of each label, given no information about the features.
 - Sometimes called the *prior*.
 - Example: What’s the probability that any given email is spam before looking at its content?
- $P(F_i|Y)$ = One table per feature. Probability of a feature, given the label.
 - Example: What’s the chance the word “FREE” appears given that the email is spam?





A Bayes Net. Each node represents a random variable. An arrow from $A \rightarrow B$ means that A influences B (i.e., B 's probability depends on A).

Domain 1: Naïve Bayes for Digits

■ Simple digit recognition version:

- One feature (variable) $F_{i,j}$ for each grid position $\langle i,j \rangle$
- Feature values are on / off, based on whether intensity is more or less than 0.5 in underlying image
- Each input maps to a feature vector, e.g.
 $F_{0,0}=0$ $F_{0,1}=0$ $F_{0,2}=1$ $F_{0,3}=1$ $F_{0,4}=1$... $F_{7,7}=0$
- Here: lots of features, each is binary valued
 - n : # of features. $n = 8 \times 8 = 64$ pixels.
 - $|F|$: # of values a single feature can take. $|F| = 2$ (0 or 1).
 - $|Y|$: # of classes. $|Y|=10$.



■ What are the parameters to learn of this model?

- $P(Y = y)$ — the prior probability of each digit class (how often each digit appears). Total # of prior probabilities to learn $= |Y|=10$.
- $P(F_{i,j} \mid Y = y)$ — the likelihood: for each class and each pixel, the probability that the pixel is on (1) or off (0). Total # of conditional probabilities to learn $= |Y| \times n \times |F| = 10 \times 64 \times 2 = 1280$.


Domain 1: Naïve Bayes for Digits: Parameters

$P(Y)$

1	0.1
2	0.1
3	0.1
4	0.1
5	0.1
6	0.1
7	0.1
8	0.1
9	0.1
0	0.1

$P(F_{3,1} = \text{on} \mid Y)$

1	0.01
2	0.05
3	0.05
4	0.30
5	0.80
6	0.90
7	0.05
8	0.60
9	0.50
0	0.80




$P(F_{5,5} = \text{on} \mid Y)$

1	0.05
2	0.01
3	0.90
4	0.80
5	0.90
6	0.90
7	0.25
8	0.85
9	0.60
0	0.80

Domain 2: Naïve Bayes for Text

■ Bag-of-words Naïve Bayes:

- Features: W_i is the word at position i
- As before: predict label conditioned on feature variables (spam vs. ham)
- As before: assume features are conditionally independent given label
- New: each W_i is identically distributed, so $P(W_1 | Y) = P(W_2 | Y) = \dots$
 - We don't care **where** a word appears, just **whether** it appears.



■ “Tied” distributions and bag-of-words

- Usually, each feature (word position) would get its own conditional probability distribution $P(W_i | Y)$
- In a bag-of-words model
 - Each position is identically distributed.
 - All positions share the same (i.e. Tied) conditional probs $P(W | Y)$.
 - Pro: simpler and more efficient. Con: the model ignores word order.
- Called “bag-of-words” because model is insensitive to word order or reordering

free our offer try please

please try our free offer

Domain 2: Naïve Bayes for Text: Parameters

■ What are the parameters?

$P(Y)$

ham	:	0.67
spam	:	0.33

If your training data has 300 emails,
with 200 ham and 100 spam:

$$P(\text{ham}) = 200/300 = 0.67$$

$$P(\text{spam}) = 100/300 = 0.33$$

$P(W|\text{spam})$

the	:	0.0156
to	:	0.0153
and	:	0.0115
of	:	0.0095
you	:	0.0093
a	:	0.0086
with	:	0.0080
offer	:	0.0035
...		

$P(W|\text{ham})$

the	:	0.0210
to	:	0.0133
of	:	0.0119
2002	:	0.0110
with	:	0.0108
from	:	0.0107
and	:	0.0105
offer	:	0.00067
...		

- n = # of words.
- $|\mathbf{F}| = 2$ (is this word in the email or not).
- $|\mathbf{Y}| = 2$ (ham or spam).

$$P(W_i|y) = \frac{\text{count of word } W_i \text{ in class } y + 1}{\text{total words in class } y + V}$$

W_i : a word in the vocabulary

y : class label (spam or ham)

V : total number of unique words in the vocabulary (used for smoothing)


The **+1** and **+V** come from **Laplace smoothing**, which prevents zero probabilities for unseen words.

Naïve Bayes Model

- In general, the joint probability (posterior) of the label and all features in Naïve Bayes model is:

$$P(Y, F_1 \dots F_n) = P(Y) \prod_i P(F_i|Y)$$

|Y| parameters
n x |F| x |Y| parameters



- We only have to specify how each feature depends on the class
- Total number of parameters is *linear* in n
 - Without the Naïve assumption (then the features can be dependent on each other), we need $|Y| \times |F|^n$ values
- Model is very simplistic, but often works anyway

Inference for Naïve Bayes

- Goal: compute posterior distribution over label variable Y

- Step 1: get joint probability of label and evidence for each label

$$P(Y, f_1 \dots f_n) = \begin{bmatrix} P(y_1, f_1 \dots f_n) \\ P(y_2, f_1 \dots f_n) \\ \vdots \\ P(y_k, f_1 \dots f_n) \end{bmatrix} \Rightarrow \begin{bmatrix} P(y_1) \prod_i P(f_i|y_1) \\ P(y_2) \prod_i P(f_i|y_2) \\ \vdots \\ P(y_k) \prod_i P(f_i|y_k) \end{bmatrix}$$

+ ↶

- Step 2: sum to get probability of evidence

- Step 3: normalize by dividing Step 1 by Step 2

↓

$$P(Y|f_1 \dots f_n)$$

Example: Spam Filtering

■ **Model:** $P(Y, W_1 \dots W_n) = P(Y) \prod_i P(W_i|Y)$

■ **Parameters:**

$P(Y)$	$P(W \text{spam})$	$P(W \text{ham})$
ham : 0.66	the : 0.0156	the : 0.0210
spam: 0.33	to : 0.0153	to : 0.0133
	you : 0.0093	

“you” is not listed: assume a small value (e.g., **Laplace smoothing** gives ~ 0.0001)

Suppose we have a short email with three words: **"the", "to", "you"**

Step 1 Compute the joint probabilities:

$$\begin{aligned} P(\text{spam}, \text{"the"}, \text{"to"}, \text{"you"}) &= P(\text{spam}) \times P(\text{"the"} | \text{spam}) \times P(\text{"to"} | \text{spam}) \times P(\text{"you"} | \text{spam}) \\ &= 0.33 \times 0.0156 \times 0.0153 \times 0.0093 \approx 7.34 \times 10^{-7} \end{aligned}$$

$$\begin{aligned} P(\text{ham}, \text{"the"}, \text{"to"}, \text{"you"}) &= P(\text{ham}) \times P(\text{"the"} | \text{ham}) \times P(\text{"to"} | \text{ham}) \times P(\text{"you"} | \text{ham}) \\ &= 0.66 \times 0.0210 \times 0.0133 \times 0.0001 \approx 1.84 \times 10^{-7} \end{aligned}$$

Step 2 Compute total probability of evidence (denominator)

$$P(\text{"the"}, \text{"to"}, \text{"you"}) = 7.34 \times 10^{-7} + 1.84 \times 10^{-7} = 9.18 \times 10^{-7}$$

Step 3 Normalize to get **posterior**

$$P(\text{spam} | \text{"the"}, \text{"to"}, \text{"you"}) = \frac{7.34 \times 10^{-7}}{9.18 \times 10^{-7}} \approx 0.799$$


$$P(\text{ham} | \text{"the"}, \text{"to"}, \text{"you"}) = \frac{1.84 \times 10^{-7}}{9.18 \times 10^{-7}} \approx 0.201$$

Final Classification → **spam**, since it has a higher posterior probability ($\approx 80\%$).

General Naïve Bayes


- What do we need in order to use Naïve Bayes?
 - Inference method (we just saw this part)
 - Start with a bunch of probabilities: $P(Y)$ and the $P(F_i|Y)$ tables
 - Use standard inference to compute $P(Y|F_1...F_n)$
 - Nothing new here
 - Estimates of local conditional probability tables
 - $P(Y)$, the prior over labels
 - $P(F_i|Y)$ for each feature (evidence variable)
 - These probabilities are collectively called the *parameters* of the model and denoted by θ
 - They typically come from training data counts


Parameter Estimation



Parameter Estimation

- Estimating the distribution of a random variable
- *Elicitation*: ask a human (why is this hard?)
- *Empirically*: use training data (learning!)
 - Example: The parameter θ is the true fraction of red beans in the jar. You don't know θ but would like to estimate it.
 - Collecting training data: You randomly pull out 3 beans:


 - Estimating θ using counts, you guess $2/3$ of beans in the jar are red.
 - Can we mathematically show that using counts is the “right” way to estimate θ ?



Parameter Estimation with Maximum Likelihood

- θ is the true fraction of red beans in the jar (i.e. $P(\text{red} \mid \theta) = \theta$)
- Can we mathematically show that using counts is the “right” way to estimate θ ?
- **Maximum likelihood estimation:** Choose the θ value that maximizes the probability of the observation
 - In other words, choose the θ value that maximizes $P(\text{observation} \mid \theta)$
 - For our problem:

$$P(\text{observation} \mid \theta)$$

$$= P(\text{randomly selected 2 red and 1 blue} \mid \theta \text{ of beans are red})$$

$$= P(\text{red} \mid \theta) * P(\text{red} \mid \theta) * P(\text{blue} \mid \theta)$$

$$= \theta^2 (1 - \theta)$$



- We want to compute:

$$\underset{\theta}{\operatorname{argmax}} \theta^2 (1 - \theta)$$

Parameter Estimation with Maximum Likelihood

- We want to compute:

$$\underset{\theta}{\operatorname{argmax}} \theta^2 (1 - \theta)$$



- Set derivative to 0, and solve!

- Common issue: The likelihood (expression we're maxing) is the product of a lot of probabilities. This can lead to complicated derivatives.
- Solution: Maximize the log-likelihood instead (this will turn products to sums), which are easier to differentiate.
- Useful fact:

$$\underset{\theta}{\operatorname{argmax}} f(\theta) = \underset{\theta}{\operatorname{argmax}} \ln f(\theta)$$

Parameter Estimation with Maximum Likelihood

$$\operatorname{argmax}_{\theta} \theta^2(1 - \theta)$$

Find θ that maximizes likelihood

$$= \operatorname{argmax}_{\theta} \ln(\theta^2(1 - \theta))$$

Find θ that maximizes log-likelihood (will be the same θ)

$$\frac{d}{d\theta} \ln(\theta^2(1 - \theta)) = 0$$

Set derivative to 0

$$\frac{d}{d\theta} [\ln(\theta^2) + \ln(1 - \theta)] = 0$$

Logarithm rule: products become sums

$$\frac{d}{d\theta} [2 \ln(\theta) + \ln(1 - \theta)] = 0$$

Logarithm rule: exponentiation becomes multiplication

$$\frac{d}{d\theta} 2 \ln(\theta) + \frac{d}{d\theta} \ln(1 - \theta) = 0$$

Now we can derive each term of the original product separately

$$\frac{2}{\theta} - \frac{1}{1 - \theta} = 0$$

Reminder: Derivative of $\ln(\theta)$ is $1/\theta$

$$\theta = \frac{2}{3}$$

Use algebra to solve for θ . If we used arbitrary red and blue counts r and b instead of $r=2$ and $b=1$, we'd get $\theta = r / (r+b)$, the count estimate.

Parameter Estimation with Maximum Likelihood (General Case)

■ Model:

X	red	blue
$P(X \theta)$	θ	$1 - \theta$

■ Data: draw N balls, N_r come up red and N_b come up blue

- Dataset $D = x_1, \dots, x_N$ of N ball draws

$$P(D|\theta) = \prod_i P(x_i|\theta) = \theta^{N_r} \cdot (1 - \theta)^{N_b}$$



■ Maximum Likelihood Estimation: find θ that maximizes $P(D|\theta)$:

$$\hat{\theta} = \underset{\theta}{\operatorname{argmax}} P(D|\theta) = \underset{\theta}{\operatorname{argmax}} \boxed{\log P(D|\theta)} \leftarrow N_r \log(\theta) + N_b \log(1 - \theta)$$

Take derivative and set to 0:

$$\begin{aligned} \frac{d}{d\theta} \log P(D|\theta) &= \frac{d}{d\theta} [N_r \log \theta + N_b \log(1 - \theta)] = N_r \frac{d}{d\theta} \log \theta + N_b \frac{d}{d\theta} \log(1 - \theta) \\ &= \frac{N_r}{\theta} + \frac{N_b}{1 - \theta} \cdot (-1) = N_r(1 - \theta) - N_b \theta = N_r - \theta(N_r + N_b) = 0 \rightarrow \hat{\theta} = \frac{N_r}{N_r + N_b} = \frac{\# \text{ of red balls}}{\text{total \# of balls}} \end{aligned}$$


Parameter Estimation with Maximum Likelihood

- Collectively name all model parameters (i.e. probability tables) as θ
- **Maximum Likelihood Estimation:** find θ that maximizes $P(\text{Data}|\theta)$
 - In practice, maximize $\log P$ instead because computation is easier
 - To solve, take derivative and set to 0
- For Naïve Bayes maximum likelihood estimates of prob. tables are:

$$P(y) = \frac{\text{\# of occurrences of class } y}{\text{total \# of observations}} \quad P(f | y) = \frac{\text{\# of occurrences of feature } f \text{ and class } y}{\text{total \# of occurrences of class } y}$$

- Need to be careful though ... let's see what can go wrong..



What is the best way to fit this data?



Empirical Risk (training error) Minimization

- How should we evaluate the quality of our model?
- Empirical risk minimization
 - Basic principle of machine learning
 - We want the model (classifier, etc) that does best on the true test distribution
 - Don't know the true distribution so pick the best model on our actual training set
 - Finding “the best” model on the training set is phrased as an optimization problem
- Main worry: overfitting to the training set
 - Better with more training data (less sampling variance)
 - Better if we limit the complexity of our hypotheses (regularization and/or small hypothesis spaces)
- Another worry: our training distribution doesn't match true distribution

Underfitting and Overfitting



Example: Overfitting

$P(\text{features}, C = 2)$


$P(C = 2) = 0.1$

$P(\text{on}|C = 2) = 0.8$

$P(\text{on}|C = 2) = 0.1$

$P(\text{off}|C = 2) = 0.1$

$P(\text{on}|C = 2) = 0.01$



$P(\text{features}, C = 3)$

$P(C = 3) = 0.1$


$P(\text{on}|C = 3) = 0.8$

$P(\text{on}|C = 3) = 0.9$

$P(\text{off}|C = 3) = 0.7$

$P(\text{on}|C = 3) = 0.0$

2 wins!!



Generalization and Overfitting

- Relative frequency parameters will **overfit** the training data!


- Just because we never saw a 3 with pixel (15,15) on during training doesn't mean we won't see it at test time
- Unlikely that every occurrence of "minute" is 100% spam
- Unlikely that every occurrence of "seriously" is 100% ham
- What about all the words that don't occur in the training set at all?
- In general, we can't go around giving unseen events zero probability

- As an extreme case, imagine using the entire email as the only feature

- Would get the training data perfect (if deterministic labeling)
- Wouldn't *generalize* at all
- Just making the bag-of-words assumption gives us some generalization, but isn't enough

- To generalize better: we need to **smooth/regularize** the estimates

Smoothing



Laplace Smoothing

■ Laplace's estimate:


- Pretend you saw every outcome once more than you actually did

$$P_{LAP}(x) = \frac{c(x) + 1}{\sum_x [c(x) + 1]}$$

Add 1 to each count as if we had seen each outcome once more than we did

$$= \frac{c(x) + 1}{N + |X|}$$

- $c(x)$: count of x
- N : total observations
- $|X|$: number of possible outcomes (e.g., red, blue $\rightarrow |X| = 2$)



$$P(\text{red}) = 2/3 \approx 0.67$$

$$P(\text{blue}) = 1/3 \approx 0.33$$

$$P_{LAP}(\text{red}) = (2+1)/(3+2) = 3/5 = 0.6$$

$$P_{LAP}(\text{blue}) = (1+1)/(3+2) = 2/5 = 0.4$$

Generalized Laplace Smoothing

■ Laplace with strength k:

- Pretend you saw every outcome k extra times

$$P_{LAP,k}(x) = \frac{c(x) + k}{N + k|X|}$$


- k is the **strength** of the prior: how strongly your prior belief (before seeing the data) influences the resulting probability estimate.
 - k = 0: You are relying **entirely** on the observed data, no smoothing → back to MLE
 - k = 1: basic Laplace
 - k > 1: stronger smoothing — useful when you have less data or expect more unseen outcomes

Generalized Laplace Smoothing

■ Laplace for conditionals:

- Smooth each condition independently:

$$P_{LAP,k}(x|y) = \frac{c(x, y) + k}{c(y) + k|X|}$$



Suppose we are trying to estimate the probability of drawing a red or blue bean given some condition y . For example, y is the day we drew the beans.

We've observed:

- $c(\text{red}, y) = 2$
- $c(\text{blue}, y) = 1$
- $c(y) = 3$ (total draws with condition y)

What's $P(\text{green} | y)$? 0! zero probability for unseen outcomes = bad for generalization!

Laplace Smoothing with $k = 1$ per condition:

- | | |
|----------------------------|---|
| • $c(\text{red}, y) = 2$ | $P(\text{red} y) = (2+1)/(3+3)=3/6=0.5$ |
| • $c(\text{blue}, y) = 1$ | $P(\text{blue} y) = (1+1)/(3+3)=2/6 \approx 0.33$ |
| • $c(\text{green}, y) = 0$ | $P(\text{green} y) = (0+1)/(3+3)=1/6=0.167$ |
| • $c(y) = 3$ | |
| • $ X =3$ | |

Naïve Bayes: No Smoothing (Overfitting)

- *Relative probabilities (odds ratios):*


$$\frac{P(W|\text{ham})}{P(W|\text{spam})}$$

south-west	:	inf
nation	:	inf
morally	:	inf
nicely	:	inf
extent	:	inf
seriously	:	inf
...		

$$\frac{P(W|\text{spam})}{P(W|\text{ham})}$$

screens	:	inf
minute	:	inf
guaranteed	:	inf
\$205.00	:	inf
delivery	:	inf
signature	:	inf
...		

What went wrong here?



Naïve Bayes: With Smoothing


- For real classification problems, smoothing is critical
- New odds ratios:

$$\frac{P(W|\text{ham})}{P(W|\text{spam})}$$

helvetica	:	11.4
seems	:	10.8
group	:	10.2
ago	:	8.4
areas	:	8.3
...		

$$\frac{P(W|\text{spam})}{P(W|\text{ham})}$$


verdana	:	28.8
Credit	:	28.4
ORDER	:	27.2
	:	26.9
money	:	26.5
...		



Do these make more sense?

Example: Naïve Bayes for Spam Filter

- Step 1: Select a ML algorithm. We choose to model the problem with Naïve Bayes.
- Step 2: Choose features to use.



Y: The label (spam or ham)	
Y	P(Y)
ham	?
spam	?

F ₁ : A feature (do I know the sender?)		
F ₁	Y	P(F ₁ Y)
yes	ham	?
no	ham	?
yes	spam	?
no	spam	?

F ₂ : Another feature (# of occurrences of FREE)		
F ₂	Y	P(F ₂ Y)
0	ham	?
1	ham	?
2	ham	?
0	spam	?
1	spam	?
2	spam	?

Example: Naïve Bayes for Spam Filter

- Step 3: Training: Use training data to fill in the probability tables.

F ₂ : # of occurrences of FREE		
F ₂	Y	P(F ₂ Y)
0	ham	0.5
1	ham	0.5
2	ham	0.0
0	spam	0.25
1	spam	0.50
2	spam	0.25

Training Data		
#	Email Text	Label
1	Attached is my portfolio.	ham
2	Are you free for a meeting tomorrow?	ham
3	Free unlimited credit cards!!!!	spam
4	Mail \$10,000 check to this address	spam
5	Sign up now for 1 free Bitcoin	spam
6	Free money free money	spam


Row 4: $P(F_2=0 \mid Y=\text{spam}) = 0.25$ because 1 out of 4 spam emails contains “free” 0 times.

Row 5: $P(F_2=1 \mid Y=\text{spam}) = 0.50$ because 2 out of 4 spam emails contains “free” 1 time.

Row 6: $P(F_2=2 \mid Y=\text{spam}) = 0.25$ because 1 out of 4 spam emails contains “free” 2 times.

Example: Naïve Bayes for Spam Filter

- Model trained on a larger dataset:




Y: The label (spam or ham)	
Y	P(Y)
ham	0.6
spam	0.4

F ₁ : A feature (do I know the sender?)		
F ₁	Y	P(F ₁ Y)
yes	ham	0.7
no	ham	0.3
yes	spam	0.1
no	spam	0.9

F ₂ : Another feature (# of occurrences of FREE)		
F ₂	Y	P(F ₂ Y)
0	ham	0.85
1	ham	0.07
2	ham	0.08
0	spam	0.75
1	spam	0.12
2	spam	0.13

Example: Naïve Bayes for Spam Filter

- Step 4: Classification
- Suppose you want to label this email from a known sender:
“**Free** food in Soda 430 today”
- Step 4.1: Feature extraction:
 - F_1 = yes, known sender
 - F_2 = 1 occurrence of “free”



Example: Naïve Bayes for Spam Filter

■ Step 4.2: Inference

■ Instantiate features (evidence):

- $F_1 = \text{yes}$
- $F_2 = 1$

■ Compute joint probabilities:


- $P(Y = \text{spam}, F_1 = \text{yes}, F_2 = 1) = P(Y = \text{spam}) P(F_1 = \text{yes} \mid \text{spam}) P(F_2 = 1 \mid \text{spam})$
 $= 0.4 * 0.1 * 0.12 = 0.0048$
- $P(Y = \text{ham}, F_1 = \text{yes}, F_2 = 1) = P(Y = \text{ham}) P(F_1 = \text{yes} \mid \text{ham}) P(F_2 = 1 \mid \text{ham})$
 $= 0.6 * 0.7 * 0.07 = 0.0294$

■ Normalize:

- $P(Y = \text{spam} \mid F_1 = \text{yes}, F_2 = 1) = 0.0048 / (0.0048 + 0.0294) = 0.14$
- $P(Y = \text{ham} \mid F_1 = \text{yes}, F_2 = 1) = 0.0294 / (0.0048 + 0.0294) = 0.86$

■ Classification result:

- 14% chance the email is spam. 86% chance it's ham.
- Or, if you don't need probabilities, note that $0.0294 > 0.0048$ and guess ham.



Y: The label (spam or ham)	
Y	P(Y)
ham	0.6
spam	0.4

F ₁ : do I know the sender?		
F ₁	Y	P(F ₁ Y)
yes	ham	0.7
no	ham	0.3
yes	spam	0.1
no	spam	0.9

F ₂ : # of occurrences of FREE		
F ₂	Y	P(F ₂ Y)
0	ham	0.85
1	ham	0.07
2	ham	0.08
0	spam	0.75
1	spam	0.12
2	spam	0.13

What we did today

- Saw our first machine learning algorithm: Naïve Bayes
 - Model is a Bayes Net where features are independent given class label
 - Classification is just inference in Bayes Nets
 - Learning is just counting feature occurrences in training data
- Saw *Maximum Likelihood* as a principled way to estimate parameters
 - Maximize probability of the data given model parameters
 - For Naïve Bayes, we solved maximization problem analytically
- Saw that fitting training data too well can cause overfitting and we can smooth it